

Detecting Fake News with Python and Machine Learning

```
In [1]: ▶ import numpy as np
import pandas as pd
import itertools
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import PassiveAggressiveClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
In [2]: ▶ #Read the data
df=pd.read_csv('news.csv')
#Get shape and head
df.shape
df.head()
```

Out[2]:

	Unnamed: 0		title	text	label
0	8476	You Can Smell Hillary's Fear	Daniel Greenfield, a Shillman Journalism Fello...		FAKE
1	10294	Watch The Exact Moment Paul Ryan Committed Pol...	Google Pinterest Digg Linkedin Reddit Stumbleu...		FAKE
2	3608	Kerry to go to Paris in gesture of sympathy	U.S. Secretary of State John F. Kerry said Mon...		REAL
3	10142	Bernie supporters on Twitter erupt in anger ag...	— Kaydee King (@KaydeeKing) November 9, 2016 T...		FAKE
4	875	The Battle of New York: Why This Primary Matters	It's primary day in New York and front-runners...		REAL

```
In [3]: ▶ #DataFlair - Get the labels
labels=df.label
labels.head()
```

Out[3]:

0	FAKE
1	FAKE
2	REAL
3	FAKE
4	REAL

Name: label, dtype: object

```
In [4]: ▶ #DataFlair - Split the dataset
x_train,x_test,y_train,y_test=train_test_split(df['text'], labels, test_si
```

```
In [5]: #DataFlair - Initialize a TfidfVectorizer  
tfidf_vectorizer=TfidfVectorizer(stop_words='english', max_df=0.7)  
#DataFlair - Fit and transform train set, transform test set  
tfidf_train=tfidf_vectorizer.fit_transform(x_train)  
tfidf_test=tfidf_vectorizer.transform(x_test)
```

```
In [6]: #DataFlair - Initialize a PassiveAggressiveClassifier  
pac=PassiveAggressiveClassifier(max_iter=50)  
pac.fit(tfidf_train,y_train)  
#DataFlair - Predict on the test set and calculate accuracy  
y_pred=pac.predict(tfidf_test)  
score=accuracy_score(y_test,y_pred)  
print(f'Accuracy: {round(score*100,2)}%')
```

Accuracy: 92.74%

```
In [7]: #DataFlair - Build confusion matrix  
confusion_matrix(y_test,y_pred, labels=['FAKE', 'REAL'])
```

```
Out[7]: array([[586,  52],  
              [ 40, 589]], dtype=int64)
```

So with this model, we have 589 true positives, 587 true negatives, 42 false positives, and 49 false negatives.

```
In [ ]: #
```